

UNIVERSITY OF SCIENCE AND TECHNOLOGY OF HANOI
DEPARTMENT OF INFORMATION AND COMMUNICATION TECHNOLOGY



MASTER THESIS

By

Dao Xuan Quy - 2440053

Title:

ROAD SEGMENTATION FOR AUTOCAR

Supervisor: Dr. Nghiem Thi Phuong

Hanoi, August – 2026

To whom it may concern,

I, Nghiem Thi Phuong, certify that the internship report of Mr. Dao Xuan Quy is qualified to be presented in the Internship Jury 2025-2026.

Hanoi, August 2026

Supervisor's signature

Nghiem Thi Phuong

Contents

Content	III
List of Figures	IV
List of Table	V
List of Abbreviations	V
Declaration	VI
Acknowledgements	VII
Abstract	VIII
1 Introduction	1
1.1 Context and motivation	1
1.1.1 Context	1
1.2 Related Work	3
1.2.1 Problem Statement	6
1.2.2 Motivation	7
1.3 Thesis Objective	7
1.4 Thesis Structure	8

2	Background	9
2.1	Kolmogorov-Arnold Networks	9
2.1.1	Original KAN Implementation with B-Splines	11
2.1.2	Optimized KAN networks	11
2.2	U-Net Segmentation Architecture	15
2.3	Binary Neural Networks	16
2.3.1	Straight-Through Estimator (STE)	17
2.3.2	Key BNN Techniques	17
2.3.3	Binarized Kolmogorov-Arnold Networks Convolutional layer (BiKA_Conv2d)	18
2.4	Loss Functions for Segmentation	19
2.4.1	Cross-Entropy Loss	19
2.4.2	Dice Loss	20
2.4.3	Combined Cross-Entropy Dice Loss	20
2.5	BDD100K Dataset	20
3	Methodology	22
3.1	UKAN: U-Shaped KAN Segmentation Network	22
3.1.1	Overall Architecture	22
3.1.2	KAN Encoder and Bottleneck	24
3.1.3	KANBlock Design	24
3.1.4	KAN Variant Layer Integration in UKAN	25
3.1.5	Decoder	26
3.2	BiKASegNet: Binarized KAN Segmentation Network	27
3.2.1	Adapting BiKA Layers for Semantic Road Segmentation	28
3.2.2	BiKAConvBlock Design	31

3.3	Optimisation and Training Strategies	31
3.3.1	Custom CUDA Kernel Optimisations for BiKA	31
3.3.2	Data Augmentation	34
4	Evaluation and Results	35
4.1	Experimental Setup	35
4.2	Evulation Metrics	36
4.3	Results and Deiscussion	37
4.3.1	Inference Throughput	38
5	Conclusions and Future Works	40
5.1	Conclusions	40
5.2	Future Works	41
	Bibliography	42

List of Figures

2.1	MLP vs KAN [1].	10
2.2	Standard U-Net encoder-decoder architecture with skip connections [2].	15
2.3	Sample 2×3 grid from the BDD100K dataset: top row shows RGB driving scene images, bottom row shows the corresponding pixel-level ground truth semantic segmentation color labels [3].	21
3.1	UKAN Architecture: Symmetric U-shaped encoder-decoder network transitioning from CNN stages to KAN blocks at lower resolutions, connected by horizontal skip connections. .	23
3.2	BiKASegNet Architecture	27

List of Tables

3.1	Side-by-side comparison between 16-channel baseline and 32-channel BiKASegNet configurations.	29
4.1	Segmentation accuracy comparison across architectures on BDD100K road segmentation.	37
4.2	Model efficiency comparison: parameter count, storage size, inference throughput, and compression ratio.	38
4.3	Per-component inference latency comparison between UKAN and BiKASegNet.	38
4.4	Ablation study on knowledge distillation and boundary loss for BiKASegNet.	39

Declaration

I hereby, Dao Xuan Quy, declare that my thesis represents my own work after doing an internship at USTH ICTLAB, University of Science and Technology of Hanoi.

I have read the University's internship guidelines. In the case of plagiarism appearing in my thesis, I accept responsibility for the conduct of the procedures in accordance with the University's Committee.

Hanoi, August 2026

Signature

Dao Xuan Quy

Acknowledgements

I would like to express my sincere gratitude to Dr. Nghiem Thi Phuong and A.Prof. Tran Giang Son for they exceptional guidance and support throughout this project. I am also grateful to USTH for providing me with the computational resources needed to complete this research, and to the open-source communities behind the KAN, BiKA and U-KAN projects for providing foundational implementations.

Abstract

Semantic segmentation is a critical component of autonomous driving perception, requiring both high accuracy and low latency for real-time deployment. Traditional segmentation networks based on U-Net [2], and Transformer architectures rely on Multi-Layer Perceptrons (MLPs) with fixed activation functions, limiting their representational flexibility. On the other hand, lightweight models often suffer from accuracy loss when handling more complicated classes. In this thesis, we study and implement two novel approaches that integrate Kolmogorov-Arnold Networks (KANs) into segmentation architectures for road segmentation. The main contributions of this thesis include:

- Optimizing and Evaluating U-KAN (UKAN), a U-Net style segmentation network that replaces standard MLP layers with KAN layers using learnable activation functions, including multiple KAN variants: FasterKAN (RSWAf basis) , ReLU-KAN, HardSwish-KAN, PWLO-KAN, and TeLU-KAN.
- Implementing and evaluating BiKA (Binarized KAN), a novel architecture that combines the KAN philosophy of per-connection learnable activations with Binary Neural Network (BNN) techniques, achieving extreme compression through 1-bit binarized weights and activations.
- Improving BiKA [4] GPU performance by apply GPU optimization methods to the official implementation
- Benchmarking both architectures against standard baselines in terms of segmentation accuracy (mIoU, Dice), model size, and inference

throughput with two distinct problems: Segment Everything (20 classes) and Segment the drivable area only (2 classes)

The experimental results demonstrate that KAN-based architectures successfully achieve a highly accurate segmentation results for multi-class segmentation with UKAN architectures, and retain acceptable accuracy while perform extremely fast with BiKA implementations.

Keywords: Kolmogorov-Arnold Networks, Semantic Segmentation, Binary Neural Networks, Road Segmentation, U-Net, Edge Deployment.

Chapter 1

Introduction

1.1 Context and motivation

1.1.1 Context

Autonomous driving car, for a long time ago have been one of human biggest dreams. Writers and filmmakers have tried different concepts of self-driving car, come-along with a modern far futures. The very first Auto-driven car came from the 80s TV series called Knight-Rider, which was a car that can drive itself and talk to the driver. In the 90s, the movie Total Recall introduced a self-driving taxi that can take passengers to their destinations without any human intervention. In the 2000s, the movie I, Robot featured a self-driving car that can navigate through traffic and avoid obstacles. And we cannot forget the absolute distopian Nigh City of Cyberpunk, packed with autonomous cars, in which the author state the constrast between the rapid rise in technology and the downfall of humanity. Another type of autonomous vehicle comes from the Galaxy far far away of Star War that symbolize the ambition of humanity to "make life easier!".

With latest advance in technology, these cars are not blueprints anymore but a fully functional vehicle. Right in the late 90s, **1977 Tsukuba Passenger Car:** - a modified vehicle built by Japan's Tsukuba Mechanical

Engineering Laborator, became the world's first machine-vision-guided semi-autonomous car. The system functioned using two vehicle-mounted optical cameras that processed live visual input using early analog machine vision algorithms to follow white street markers. In the end, it could reach 30 km/h on dedicated marked test tracks. Six year later **1986 Carnegie Mellon University Navlab 1:** was constructed by Carnegie Mellon University, marking the first vehicle driven autonomously using artificial neural networks. Researchers developed it to transition from rule-based computer vision toward adaptive machine learning systems capable of handling complex road geometries. It processed camera images using the ALVINN neural network architecture running on five onboard computer racks powered by a 5 kW generator. Its top performance reached an autonomous driving speed of 32 km/h (20 mph) on campus roads. With the development in Deep learning neural networks, autotnomous cas now combining camera with other sensor to better environment visualization. With the safeness increase, **2015 Tesla Autopilot (HW1):** Tesla Autopilot (HW1) was a mass-market Advanced Driver Assistance System (ADAS) introduced by Tesla, Inc. for commercial electric vehicles. The latest **2026 Waymo Driver (Level 4)** driverless fleet operating commercial robo-taxi services. It provide fully automated taxi services. The system relies on custom Level 4 sensor suites including high-resolution LiDAR, radar, 360-degree cameras, and deep AI foundation models trained on billions of real-world and simulated miles. Its top performance includes logging over 220.6 million commercial rider-only miles while achieving an 84% reduction in injury-causing crashes compared to human drivers.

Although many autonomous car driving system have been created, they share the same mechanics: use the camera as an "eyes" of the car, capture the environment to identify other cars, obstacles, pedestrians and drivable lane with a semantic segmentation neural networks - the task of assigning a class label to every pixel in an image. Accurate road segmentation enables autonomous vehicles to identify drivable surfaces, lane boundaries, sidewalks, and other environment objects in real time. The evolution of seg-

mentation architectures has progressed from early fully convolutional networks (FCN) [5] to the U-Net encoder-decoder architecture [2], and more recently to Transformer-based models such as SegFormer [6], Swin-UNet [7] and Segment Anything [8].

1.2 Related Work

The evolution of semantic segmentation began with the use of classification Convolutional Neural Networks (CNNs) for dense pixel-level prediction. Long et al. [5] introduced Fully Convolutional Networks (FCN), replacing fully connected layers with 1×1 convolutions to preserve spatial coordinates and adapt to dynamic input shapes; by combining deep, coarse semantic features with shallow, fine spatial details via transposed convolutions (FCN-8s), FCN achieved 62.2% mean Intersection over Union (mIoU) on PASCAL VOC 2012 and 34.0% mIoU on NYUDv2 [9]. To address the loss of fine-grained spatial localization caused by repetitive pooling operations, Ronneberger et al. [2] proposed U-Net, a symmetric encoder-decoder architecture that concatenates high-resolution encoder feature maps directly to decoder stages across horizontal skip connections, establishing a dominant paradigm in biomedical and road segmentation while achieving 92.0% IoU on the ISBI Cell Tracking Challenge. Seeking to minimize memory consumption during upsampling, Badrinarayanan et al. [10] developed SegNet, which records max-pooling indices during encoding and reuses them in the decoder to perform sparse parameter-free unpooling, reaching 60.1% mIoU on CamVid and 57.0% mIoU on Cityscapes.

To resolve contextual ambiguities and capture multi-scale spatial features simultaneously, context aggregation architectures were developed. Chen et al. [11] proposed DeepLabV3+, combining Atrous Spatial Pyramid Pooling (ASPP) with an encoder-decoder structure; by employing dilated convolutions at multiple sampling rates ($r = 6, 12, 18$) to expand the receptive field without increasing parameter counts, DeepLabV3+ achieved state-of-the-art performance with 89.0% mIoU on PASCAL VOC 2012 and 82.1% mIoU on

Cityscapes. Similarly, Zhao et al. [12] introduced the Pyramid Scene Parsing Network (PSPNet) to prevent global context misclassifications—such as confusing boats on water with land vehicles—by applying a 4-level Pyramid Pooling Module ($1 \times 1, 2 \times 2, 3 \times 3, 6 \times 6$) to aggregate sub-region features across multiple scales, obtaining 80.2% mIoU on Cityscapes and 85.4% mIoU on PASCAL VOC 2012.

Deploying semantic segmentation models on autonomous vehicles requires high processing throughput under strict computational and thermal constraints. Paszke et al. [13] designed ENet, an ultra-compact real-time architecture utilizing bottleneck blocks with early spatial downsampling, factorized asymmetric 5×1 and 1×5 convolutions, and dilated kernels to maintain receptive fields with minimal parameters, reaching 58.3% mIoU at 76.8 frames per second (FPS) on an NVIDIA Titan X for 1024×512 Cityscapes images and 51.3% mIoU on CamVid. To overcome the speed-accuracy trade-off inherent in aggressive downsampling, Yu et al. [14] proposed BiSeNet (Bilateral Segmentation Network), which decouples spatial detail preservation from context extraction into a shallow Spatial Path and a deep Context Path with Attention Refinement Modules; BiSeNet achieved an impressive 105.8 FPS with 74.7% mIoU on Cityscapes (768×1536 resolution) and 68.7% mIoU on CamVid at 146.7 FPS.

Attention mechanisms have further advanced dense visual perception by capturing global, long-range dependencies across driving scenes. Xie et al. [6] introduced SegFormer, a lightweight, positional-encoding-free Transformer paired with an All-MLP decoder; by employing overlapping patch embeddings and spatial-reduction self-attention, SegFormer achieved 84.0% mIoU with MiT-B5 on Cityscapes (and 76.2% mIoU at 47.7 FPS with MiT-B0), as well as 51.8% mIoU on ADE20K. To eliminate the quadratic computational complexity $\mathcal{O}(H^2W^2)$ of global self-attention while maintaining cross-window connectivity, Cao et al. [7] proposed Swin-UNet, a pure U-shaped Transformer architecture replacing convolutional layers with Shifted Window (S-WMSA) self-attention blocks, achieving 79.13% average Dice score on the Synapse Multi-Organ dataset and 90.00% Dice on ACDC. Addressing

zero-shot class-agnostic segmentation, Kirillov et al. [8] developed Segment Anything (SAM), training a heavy ViT image encoder paired with a prompt encoder and two-way transformer mask decoder on the 1.1-billion-mask SA-1B dataset to achieve 46.5 mIoU zero-shot on COCO.

Recently, Kolmogorov-Arnold Networks (KANs) were introduced by Liu et al. [1] as a fundamental alternative to traditional Multi-Layer Perceptrons (MLPs). Grounded in the Kolmogorov-Arnold Representation Theorem [15], KANs replace fixed node activation functions (e.g., ReLU, GELU) with learnable 1D B-spline activation functions on network edges ($\phi(x) = w_b b(x) + w_s \text{spline}(x)$), enabling the network to dynamically learn activation function shapes while outperforming MLPs with up to $10\times$ fewer parameters on PDE solving and function approximation. Extending KANs to dense vision tasks, Li et al. [16] introduced U-KAN (UKAN), reformulating 1D spline KAN operations into vectorized 2D matrix multiplications to replace standard MLP/Transformer blocks within a U-Net backbone, attaining an 89.84% average Dice score on the Synapse Multi-Organ dataset and 90.12% IoU on ISIC 2018 Skin Lesion Segmentation. To resolve the slow GPU training latency caused by non-uniform B-spline grid allocations in KANs, Tsiamis [17] developed FasterKAN, replacing grid-allocated B-splines with Reflectional Switch Activation Functions (RSWAf) based on multiscale Radial Basis Functions (RBF) to achieve a $3\times$ to $5\times$ speedup on CUDA platforms.

To achieve extreme model compression and speed for edge hardware deployment, Binary Neural Networks (BNNs) constrain both weights and activations to 1-bit binary values ($\{-1, +1\}$). Hubara et al. [18] formalized 1-bit neural networks using the sign binarization function $x_b = \text{sign}(x)$ in the forward pass and the Straight-Through Estimator (STE) during backpropagation, replacing expensive 32-bit floating-point multiply-accumulate operations with bitwise XNOR logic and population count (`__popc`) hardware intrinsics, achieving 98.6% accuracy on MNIST and 89.8% on CIFAR-10. To mitigate quantization loss, Rastegari et al. [19] proposed XNOR-Net by introducing analytic real-valued channel scaling factors (α), while Liu et al. [20] introduced Bi-Real Net with FP32 real-valued identity shortcuts

($\mathbf{y} = \text{sign}(\mathbf{x}) * \mathbf{W}_b + \mathbf{x}$) to preserve continuous gradient flow, reaching Top-1 ImageNet accuracies of 51.2% and 56.4%, respectively. Further advancing BNN expressiveness, ReActNet [21] introduced learnable per-channel activation shifts via RReLU ($\text{RReLU}(x) = \text{ReLU}(x + \gamma) + \beta$), IR-Net [22] proposed Libra-PB for information-entropy-balanced weight binarization, and AdaBin [23] fitted adaptive binary scaling sets, enabling ReActNet to reach 69.4% Top-1 accuracy on ImageNet. Extending binary compression to dense prediction, Zhuang et al. [24] presented Group-Net, a structured binary network that divides binary operations into multiple parallel branches with group fusion to reconstruct continuous spatial feature maps, achieving 63.8% mIoU on Cityscapes road segmentation and 68.1% mIoU on PASCAL VOC under 1-bit quantization.

1.2.1 Problem Statement

However, all of these architectures share a fundamental design choice inherited from Multi-Layer Perceptrons (MLPs): they place fixed, pre-defined activation functions (e.g., ReLU, GELU) at the *nodes* of the network, while the *edges* (connections) carry only learnable linear weights. The Kolmogorov-Arnold Representation Theorem [15, 25] suggests an alternative: replace every linear weight with a non-linear activation functions on the *edges* themselves, allowing the network to learn both the weights and the optimal activation shape for each connection, which as a results, achieve better accuracy and interpretability. This principle forms the foundation of Kolmogorov-Arnold Networks (KANs) [1].

In the Sematic Segmentation contetxt, recent work has demonstrated the potential of KAN-based architectures in medical image segmentation [16]. On the other hand, KAN application to road scene segmentation and their compatibility with extreme model compression techniques such as Binary Neural Networks (BNNs) [20, 21] remain largely unexplored.

Deploying segmentation models on autonomous vehicles poses two competing requirements. First, the model must achieve high segmentation accuracy to

to ensures safeness to vehicles and the passengers. Second, since the car **sees** the environment throught a camera, the model itself have to be capable of processing a large amount of frames per second; therefore, require low latency, small model size, and minimal memory bandwidth.

Standard full-precision segmentation networks (32-bit floating point) achieve high accuracy but require significant computational resources. Binary Neural Networks (BNNs) offer extreme compression by quantizing both weights and activations to 1-bit representations, but they suffer from severe accuracy degradation due to the limited expressiveness of the sign function used for binarization. The central challenge addressed in this thesis is: *Can the KAN solve the accuracy loss in efficient, binarized networks, while preserving their computational efficiency advantages?*

1.2.2 Motivation

This research is motivated by the need to bridge the gap between expressive, high-accuracy segmentation models and the efficient model real-time performance. We perform two approaches: UKAN for full-precision KAN-augmented segmentation and BiKA for binarized KAN convolutions—we aim to observe how KAN integration affects the accuracy-efficiency for road segmentation.

1.3 Thesis Objective

The primary objective of this thesis is to implement, evaluate, and compare two KAN-based segmentation architectures for road segmentation problem. Specifically, we aim to:

- Adapt the U-KAN architecture for road segmentation on the BDD100K dataset, exploring multiple KAN linear variants (FasterKAN, ReLU-KAN, HardSwish-KAN, PWLO-KAN, TeLU-KAN) to assess their impact on segmentation accuracy and inference speed.

- Design and implement BiKASegNet, a binarized segmentation network that applies the KAN algorithm through per-connection binary thresholds to achieve extreme model compression while retaining acceptable accuracy.
- Benchmark both architectures against standard baseline models in terms of segmentation quality (mIoU, Dice), model size, and inference speed on CUDA platforms.

1.4 Thesis Structure

The remainder of this thesis is structured as follows:

- **Chapter 1** introduces the context, problem statement, related work, and objectives of this research.
- **Chapter 2** discusses the technical background of Kolmogorov-Arnold Networks, the U-Net segmentation, Binary Neural Networks, and the KAN variant implementations.
- **Chapter 3** explains the methodology, including the UKAN architecture, KAN layer integration, the BiKA binarized segmentation network design.
- **Chapter 4** describes the experimental setup, dataset preparation, training procedures, and performance evaluation of both architectures.
- **Chapter 5** summarizes the content of this thesis and proposes future developments of research.

Chapter 2

Background

This chapter introduces the foundational knowledge that are used to implement KAN-based segmentation networks. These include Kolmogorov-Arnold Networks (KANs), Binary Neural Networks (BNNs), and Binarized Kolmogorov-Arnold Networks (BiKA) concept which successfully combine the two architectures.

2.1 Kolmogorov-Arnold Networks

The Kolmogorov-Arnold Representation Theorem [15, 25] states that any multivariate continuous function $f : [0, 1]^n \rightarrow \mathbb{R}$ can be represented as a finite composition of continuous functions of a single variable and the binary operation of addition:

$$f(x_1, \dots, x_n) = \sum_{q=0}^{2n} \Phi_q \left(\sum_{p=1}^n \phi_{q,p}(x_p) \right) \quad (2.1)$$

where $\phi_{q,p} : [0, 1] \rightarrow \mathbb{R}$ are univariate inner functions and $\Phi_q : \mathbb{R} \rightarrow \mathbb{R}$ are outer functions. This theorem guarantees the existence of such a representation but does not prescribe how to learn it.

Kolmogorov-Arnold Networks (KANs) [1] have been proposed as a novel approach to improve traditional Multi-Layer Perceptrons (MLPs) accuracy,

inspired by the Kolmogorov-Arnold representation theorem. In a standard MLP, the network learns fixed scalar weights on its connections (edges) and applies pre-defined, fixed non-linear activation functions (such as ReLU or sigmoid) at its neurons (nodes). KANs invert this principle by placing learnable, non-linear activation functions directly on the network edges, while the nodes simply sum the incoming signals. The key different between KANs and traditional MLPs is illustrated in Figure 2.1.




Figure 2.1: MLP vs KAN [1].

In an MLP, each neuron computes:

$$y_j = \sigma \left(\sum_i w_{ij} x_i + b_j \right) \quad (2.2)$$

where σ is a fixed activation function (e.g., ReLU) and w_{ij} are learnable scalar weights. In a KAN, each node simply sums its incoming signals, and each edge applies a learnable univariate function:

$$y_j = \sum_i \phi_{ij}(x_i) \quad (2.3)$$

where $\phi_{ij} : \mathbb{R} \rightarrow \mathbb{R}$ are learnable activation functions parameterized as spline functions. This places the representational power on the edges rather than the nodes.

2.1.1 Original KAN Implementation with B-Splines

In the original KAN [1], each edge activation ϕ_{ij} is parameterized as a B-spline:

$$\phi_{ij}(x) = w_b \cdot \text{SiLU}(x) + w_s \cdot \text{spline}(x) \quad (2.4)$$

where $\text{SiLU}(x) = x \cdot \sigma(x)$ is a base activation, and $\text{spline}(x) = \sum_k c_k B_k(x)$ is a linear combination of B-spline basis functions B_k with learnable coefficients c_k . The B-spline is defined over a grid of knot points $\{t_0, t_1, \dots, t_G\}$ spanning the input domain.

As a trade-off to the extra information captured, B-spline evaluation requires iterative recursion over the grid, resulting in $\mathcal{O}(G \cdot k)$ operations per activation (where G is the grid size and k is the spline order), making the original KAN impractical for high-throughput vision tasks. This is the main problem that we address in this thesis.

2.1.2 Optimized KAN networks

To overcome the computational bottleneck of B-splines, several efficient KAN linear replacements have been proposed. Each variant maintains the KAN philosophy of learning activation functions on edges but uses faster basis functions.

FastKAN [26] reframes Kolmogorov-Arnold Networks through the lens of Radial Basis Function (RBF) networks. It replaces the computationally intensive 3rd-order B-splines with Gaussian RBF basis functions defined over a uniform grid of center points $\{g_0, g_1, \dots, g_{G-1}\}$:

$$\psi_k(x) = \exp \left(- \left(\frac{x - g_k}{\gamma} \right)^2 \right) \quad (2.5)$$

where $\gamma > 0$ controls the bandwidth of the Gaussian basis functions. Each univariate activation function on edge (i, j) is constructed as a linear com-

bination of these Gaussian RBFs alongside a base linear projection:

$$\phi_{ij}(x) = w_{\text{base}} \cdot \text{SiLU}(x) + \sum_{k=0}^{G-1} c_{k,ij} \cdot \psi_k(x) \quad (2.6)$$

where $c_{k,ij}$ are learnable expansion coefficients. By avoiding the iterative grid recursion required by B-splines, FastKAN computes basis evaluations in parallel via matrix multiplication, enabling execution speed comparable to standard MLPs while retaining KAN’s edge-centric non-linear expressiveness.

PowerMLP [27] addresses KAN’s training latency bottleneck by reformulating the edge-centric spline activation functions into non-iterative power-basis expansions. Instead of evaluating localized grid splines or RBFs, PowerMLP approximates each non-linear univariate edge function $\phi(x)$ using a polynomial power-series expansion of order K :

$$\phi(x) = \sum_{k=1}^K w_k \cdot \sigma(x)^k \quad (2.7)$$

where $\sigma(x)$ is a base non-linear activation function (such as SiLU or GELU), and w_k are learnable scalar weights associated with each power k . The layer output for an input vector $\mathbf{x} \in \mathbb{R}^{d_{\text{in}}}$ is computed via vectorized elementwise powers and linear projections:

$$y = \sum_{k=1}^K \mathbf{W}_k (\sigma(\mathbf{x})^k) \quad (2.8)$$

where $\mathbf{W}_k \in \mathbb{R}^{d_{\text{out}} \times d_{\text{in}}}$ are standard weight matrices. Because power calculations $(\mathbf{x}^k)$ are fully parallelizable and require no grid lookups or localized indexing, PowerMLP drastically reduces FLOPs and achieves up to 40× faster training speed than standard B-spline KANs, combining MLP-like execution throughput with KAN-level non-linear representational power.

FasterKAN [17] replaces B-splines with the Reflectional Switch Activation Function (RSWAf). Given a set of grid points $\{g_0, g_1, \dots, g_{G-1}\}$ uniformly

spanning $[g_{\min}, g_{\max}]$ and an inverse denominator parameter β , the basis function is:

$$r_k(x) = 1 - \tanh^2((x - g_k) \cdot \beta) \quad (2.9)$$

Each basis function r_k produces a smooth, localized bump centred at grid point g_k . The output of a FasterKAN linear layer with input dimension d_{in} and output dimension d_{out} is:

$$y = \mathbf{W}_s \cdot \text{vec} \left(\{r_k(x_i)\}_{k=0, \dots, G-1}^{i=1, \dots, d_{\text{in}}} \right) \quad (2.10)$$

where $\mathbf{W}_s \in \mathbb{R}^{d_{\text{out}} \times (d_{\text{in}} \cdot G)}$ is a learnable spline projection matrix, and $\text{vec}(\cdot)$ concatenates all basis evaluations into a single vector. The input is first normalised via LayerNorm. A custom CUDA kernel computes the forward and backward passes of the RSWAf function, achieving significant speedups over the B-spline implementation.

ReLU-KAN [28] replaces the smooth basis with a piecewise-linear function using ReLU:

$$\psi_k(x) = \text{ReLU}(x - g_k) = \max(0, x - g_k) \quad (2.11)$$

This produces a basis of shifted ramp functions. The output is computed as:

$$y = \mathbf{W} \cdot \text{vec}(\{\psi_k(x_i)\}_{k,i}) \quad (2.12)$$

ReLU-KAN achieves maximal inference speed due to the trivial cost of ReLU evaluation, but trades off smoothness for speed.

With the same policy proposed by FasterKAN [17] and ReLU-KAN [28], we decide to perform evaluation with more activation functions which have a lower computational complexity compared original KAN to find out the optimal point between efficiency and accuracy.

- HardSwish [29] is defined mathematically as:

$$\text{HardSwish}(x) = x \cdot \frac{\text{ReLU6}(x + 3)}{6} \quad (2.13)$$

It serves as a computationally efficient, piecewise-linear approxima-

tion of the smooth Swish activation function ($x \cdot \sigma(x)$). By replacing the continuous exponential function in Sigmoid with a bounded linear ramp ($\text{ReLU6}(x) = \min(\max(0, x), 6)$), HardSwish enforces piecewise linearity across the input domain. For $x \leq -3$, the output is identically 0; for $x \geq 3$, the function becomes purely linear ($f(x) = x$). In the intermediate region $[-3, 3]$, it forms a quadratic-like piecewise linear transition. This piecewise linear structure eliminates expensive transcendental operations while maintaining gradient stability for deep networks.

- The Hyperbolic Tangent Exponential Linear Unit (TeLU) [30] is formulated as:

$$\text{TeLU}(x) = x \cdot \tanh(\exp(x)) \quad (2.14)$$

TeLU is designed to combine the fast convergence of linear activations with the smooth non-linear saturation of exponential units. In the active positive domain ($x > 0$), $\exp(x)$ increases rapidly, driving $\tanh(\exp(x)) \approx 1$. Consequently, $\text{TeLU}(x) \approx x$, displaying strong positive-region linearity (identity mapping) that allows gradients to flow unattenuated. In the negative domain ($x < 0$), the smooth decay of $\tanh(\exp(x))$ prevents dying activations while bounding negative values, maintaining a seamless transition between linear and non-linear regimes.

- Piecewise Linear Optimisation (PWLO) [31] approximates non-linear continuous functions by partitioning the input domain $[g_{\min}, g_{\max}]$ into G sub-intervals $[g_k, g_{k+1})$ and assigning a localized linear affine transformation to each segment:

$$f(x) = a_k \cdot x + b_k, \quad x \in [g_k, g_{k+1}) \quad (2.15)$$

where a_k and b_k denote the slope and offset of the k -th segment. PWLO explicitly enforces local linearity within each interval. By converting complex non-linear functions into a composition of linear segments, PWLO enables high-throughput execution using simple linear primi-

tives (multiply-accumulate or shift-and-add) on edge acceleration hardware.

2.2 U-Net Segmentation Architecture

U-Net [2] established the symmetric encoder-decoder architecture with skip connections as the standard for dense prediction tasks. The encoder progressively downsamples the input through convolutional blocks and pooling operations, extracting increasingly abstract feature representations. The decoder upsamples these features back to the original resolution, with skip connections from corresponding encoder stages providing fine-grained spatial detail.

Each encoder stage consists of a double convolution block:

$$\text{ConvLayer}(x) = \text{ReLU}(\text{BN}(\text{Conv}_{3 \times 3}(\text{ReLU}(\text{BN}(\text{Conv}_{3 \times 3}(x)))))) \quad (2.16)$$

followed by 2×2 max pooling for spatial downsampling. The decoder mirrors this structure using bilinear upsampling and concatenation (or addition) with encoder skip connections, as illustrated in Figure 2.2.




Figure 2.2: Standard U-Net encoder-decoder architecture with skip connections [2].

2.3 Binary Neural Networks

Binary Neural Networks (BNNs) address the hardware constraints of deep learning models by simplify the mathematics back to Binary operations. The binarization function for weights is:

$$w_b = \text{sign}(w) = \begin{cases} +1 & \text{if } w \geq 0 \\ -1 & \text{otherwise} \end{cases} \quad (2.17)$$

and similarly for activations: $x_b = \text{sign}(x)$. This allows the multiply-accumulate operations in convolutions to be replaced by XNOR and popcount operations:

$$w_b \cdot x_b = \text{XNOR}(w_b, x_b), \quad \sum_i w_{b,i} \cdot x_{b,i} = \text{popcount}(\text{XNOR}(\mathbf{w}_b, \mathbf{x}_b)) \quad (2.18)$$

achieving up to $\sim 58\times$ theoretical speedup and $\sim 32\times$ memory reduction compared to full-precision networks [19]. In a BNN, both the network weights and the intermediate feature activations are constrained to 1-bit representations, typically $\{-1, +1\}$, utilizing the sign function. This method completely eliminates the need for expensive floating-point multiply-accumulate (MAC) operations. Instead, convolutions and dense layers can be executed using highly efficient bitwise XNOR and popcount operations, drastically reducing memory bandwidth and latency on edge devices. Despite these immense computational advantages, the extreme information loss induced by binarizing continuous variables natively limits the expressiveness of BNNs. They often suffer from accuracy compared to full-precision networks and require specialized training methodologies, such as the Straight-Through Estimator (STE) and dense skip connections, to maintain stable gradient flow during backpropagation.

2.3.1 Straight-Through Estimator (STE)

The sign function has zero gradient almost everywhere, making direct backpropagation impossible. Bengio et al. [32] proposed the Straight-Through Estimator (STE), which approximates the gradient of the sign function as:

$$\frac{\partial \mathcal{L}}{\partial w} \approx \frac{\partial \mathcal{L}}{\partial w_b} \cdot \mathbf{1}_{|w| \leq 1} \quad (2.19)$$

where the gradient is passed through unchanged within the clipping window $[-1, 1]$ and zeroed outside. This enables training of binary networks with standard gradient-based optimisers.

2.3.2 Key BNN Techniques

Several techniques have been developed to improve the accuracy of BNNs:

Bi-Real Net Skip Connections

Bi-Real Net [20] introduces full-precision (FP32) identity shortcut connections that bypass the binary convolution blocks:

$$y = \text{BinaryConv}(x) + \text{Shortcut}(x) \quad (2.20)$$

These skip connections serve as gradient highways, providing a direct path for gradient flow during backpropagation and significantly stabilizing training.

RPreLU (ReActNet)

ReActNet [21] replaces standard ReLU with RPreLU, a learnable activation with per-channel pre-shift and post-shift parameters:

$$\text{RPreLU}(x) = \text{PReLU}(x + \gamma) + \beta \quad (2.21)$$

where $\gamma \in \mathbb{R}^C$ and $\beta \in \mathbb{R}^C$ are learnable channel-wise shift parameters. This enables the binary network to learn asymmetric activation distributions, compensating for the information loss caused by binarization.

Libra-PB (Balanced Binarization)

IR-Net [22] proposed Libra Parameter Binarization (Libra-PB), which centres the weight distribution before applying the sign function:

$$w_b = \text{sign}(w - \bar{w}), \quad \bar{w} = \frac{1}{|\mathcal{F}|} \sum_{i \in \mathcal{F}} w_i \quad (2.22)$$

where $\bar{w}$ is the mean over the filter dimensions. This ensures approximately 50% of the binary weights are +1 and 50% are -1, maximizing the information entropy of the binary representation.

AdaBin (Adaptive Binary Scaling)

AdaBin [23] introduces a learnable per-channel output scaling factor α_c that modulates the binary convolution output:

$$y_c = \alpha_c \cdot \text{BinaryConv}_c(x) \quad (2.23)$$

At inference, α_c is fused into the subsequent BatchNorm layer’s affine parameters, incurring zero additional computational cost.

2.3.3 Binarized Kolmogorov-Arnold Networks Convolutional layer (BiKA_Conv2d)

The foundational building block of the Binarized Kolmogorov-Arnold Network is the BiKA_Conv2d layer. In a conventional 2D convolution, the learnable weight tensor has shape $(C_{\text{out}}, C_{\text{in}}, K_H, K_W)$ while the bias is a 1D vector of shape $(C_{\text{out}},)$ —assigning a single scalar offset per output channel.

In contrast, **BiKA_Conv2d** constructs a 4D per-connection bias tensor matching the exact spatial and channel dimensions of the weight tensor:

$$\mathbf{W} \in \mathbb{R}^{C_{\text{out}} \times C_{\text{in}} \times K_H \times K_W}, \quad \mathbf{b} \in \mathbb{R}^{C_{\text{out}} \times C_{\text{in}} \times K_H \times K_W} \quad (2.24)$$

This per-connection bias design realizes the KAN paradigm by assigning an individualized, learnable step-function threshold b_{o,c,k_h,k_w} to every single synapse. The forward mathematical operation for an output spatial location (h', w') is defined as:

$$y_{o,h',w'} = \sum_{c=1}^{C_{\text{in}}} \sum_{k_h=1}^{K_H} \sum_{k_w=1}^{K_W} \text{sign}(w_{o,c,k_h,k_w} - \bar{w}_o) \cdot \text{sign}(x_{c,h'+k_h,w'+k_w} + b_{o,c,k_h,k_w}) \quad (2.25)$$

where $\bar{w}_o = \frac{1}{C_{\text{in}} K_H K_W} \sum_{c,k_h,k_w} w_{o,c,k_h,k_w}$ is the Libra-PB weight mean. Each synapse evaluates a 1-bit step activation $\phi_{o,c,k_h,k_w}(x) = \text{sign}(x + b_{o,c,k_h,k_w})$. Following the binary convolution, a learnable AdaBin channel scale $\alpha_c \in \mathbb{R}^{C_{\text{out}} \times 1 \times 1}$ rescales the output tensor:

$$y_c \leftarrow \alpha_c \cdot y_c \quad (2.26)$$

2.4 Loss Functions for Segmentation

Semantic segmentation typically employs a combination of loss functions to handle class imbalance and optimize for region-based metrics.

2.4.1 Cross-Entropy Loss

The pixel-wise cross-entropy loss for C classes is:

$$\mathcal{L}_{\text{CE}} = -\frac{1}{|\Omega|} \sum_{i \in \Omega} \sum_{c=1}^C y_{i,c} \log(\hat{y}_{i,c}) \quad (2.27)$$

where Ω is the set of valid pixels, $y_{i,c}$ is the ground-truth one-hot label, and $\hat{y}_{i,c}$ is the predicted probability after softmax.

2.4.2 Dice Loss

The Dice loss directly optimizes the Dice coefficient (F1 score), which is robust to class imbalance:

$$\mathcal{L}_{\text{Dice}} = 1 - \frac{1}{C} \sum_{c=1}^C \frac{2 \sum_i y_{i,c} \hat{y}_{i,c} + \epsilon}{\sum_i y_{i,c} + \sum_i \hat{y}_{i,c} + \epsilon} \quad (2.28)$$

where ϵ is a smoothing constant to prevent division by zero.

2.4.3 Combined Cross-Entropy Dice Loss

The default loss function used in both UKAN and BiKA training combines both objectives:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{CE}} + \lambda_{\text{Dice}} \cdot \mathcal{L}_{\text{Dice}} \quad (2.29)$$

This jointly optimizes for per-pixel classification accuracy and region-level overlap, with the Dice component ensuring that small classes (e.g., lane markings) are not overwhelmed by the dominant background class.

2.5 BDD100K Dataset

The Berkeley DeepDrive 100K (BDD100K) dataset [3] is one of the largest and most diverse driving scene benchmarks designed for multitask learning in autonomous driving perception. It contains 100,000 video clips captured across varied geographic locations, times of day (daytime, night, dawn/dusk), and weather conditions (clear, rainy, foggy, snowy).

For semantic segmentation, BDD100K provides pixel-level ground truth annotations for 10,000 high-resolution images (1280×720), partitioned into

7,000 training, 1,000 validation, and 2,000 test images. The annotations cover 19 evaluation semantic classes (plus background), including road, sidewalk, vehicle, pedestrian, rider, traffic sign, and sky. In this work, we utilize BDD100K to benchmark our segmentation networks under both multi-class and binary road segmentation setups.


Figure 2.3 displays a 2×3 grid showcasing sample input driving images alongside their corresponding color-coded ground truth semantic segmentation masks.



(a) Scene Image 1



(b) Scene Image 2



(c) Scene Image 3



(d) Ground Truth 1



(e) Ground Truth 2



(f) Ground Truth 3

Figure 2.3: Sample 2×3 grid from the BDD100K dataset: top row shows RGB driving scene images, bottom row shows the corresponding pixel-level ground truth semantic segmentation color labels [3].

Chapter 3

Methodology

This chapter details the two KAN-based segmentation architectures: the full-precision UKAN and the binarized BiKASegNet. We describe the integration of KAN layers into the U-Net framework, the design of binarized KAN convolutions, and the knowledge distillation pipeline connecting both models.


3.1 UKAN: U-Shaped KAN Segmentation Network

UKAN adapts the U-Net architecture by replacing the standard MLP/Transformer blocks in the bottleneck and lower-resolution stages with KAN-based blocks, while retaining efficient CNN stages at higher resolutions where spatial dimensions are large.

3.1.1 Overall Architecture

The UKAN architecture follows a three-stage CNN encoder, a KAN encoder-bottleneck, and a mixed KAN-CNN decoder, as illustrated in Figure [3.1](#).

The encoder consists of three CNN stages (ConvLayer), each performing a



**Note: Green blocks indicate KAN-based stages. Blue blocks indicate standard CNN stages. Red block indicates the KAN bottleneck. Dashed lines indicate horizontal skip connections across the U-shape.*

Figure 3.1: UKAN Architecture: Symmetric U-shaped encoder-decoder network transitioning from CNN stages to KAN blocks at lower resolutions, connected by horizontal skip connections.

double 3×3 convolution with BatchNorm and ReLU, followed by 2×2 max pooling. The channel dimensions progress as $3 \rightarrow C_0/8 \rightarrow C_0/4 \rightarrow C_0$, where C_0 is the base embedding dimension (default: 64).

3.1.2 KAN Encoder and Bottleneck

At the transition from CNN to KAN stages, PatchEmbed layers convert 2D feature maps into 1D token sequences via a strided convolution projection:

$$\mathbf{T} = \text{LayerNorm}(\text{reshape}(\text{Conv}_{k \times k, s}(\mathbf{F}))) \in \mathbb{R}^{B \times N \times C} \quad (3.1)$$

where $\mathbf{F}$ is the input feature map, k is the patch size, s is the stride, and $N = H' \times W'$ is the number of tokens. Two KAN encoder stages process tokens at embedding dimensions $C_1 = 128$ and $C_2 = 256$ (default configuration).

3.1.3 KANBlock Design

Each KANBlock follows a Transformer-style pre-norm residual pattern:

$$\mathbf{T}' = \mathbf{T} + \text{DropPath}(\text{KANLayer}(\text{LayerNorm}(\mathbf{T}), H, W)) \quad (3.2)$$

The KANLayer replaces the standard MLP (or self-attention) module. It consists of three KAN linear projections interleaved with depthwise convolution blocks (DW_bn_relu) for spatial token mixing:

$$\mathbf{T}_1 = \text{DW_bn_relu}(\text{KAN_fc}_1(\mathbf{T}), H, W) \quad (3.3)$$

$$\mathbf{T}_2 = \text{DW_bn_relu}(\text{KAN_fc}_2(\mathbf{T}_1), H, W) \quad (3.4)$$

$$\mathbf{T}_3 = \text{DW_bn_relu}(\text{KAN_fc}_3(\mathbf{T}_2), H, W) \quad (3.5)$$

Each KAN_fc is an instance of the selected KAN linear variant (FasterKAN, ReLU-KAN, etc.) with grid size $G = 8$. The depthwise convolutions preserve spatial context since KAN operations act independently on the channel

dimension. Each DW_bn_relu block performs:

$$\text{DW_bn_relu}(\mathbf{T}, H, W) = \text{ReLU}(\text{BN}(\text{DWConv}_{3 \times 3}(\text{reshape}(\mathbf{T}, H, W)))) \quad (3.6)$$

3.1.4 KAN Variant Layer Integration in UKAN

While Chapter 2 described the mathematical definitions and linearity properties of the individual activation functions, this section details how HardSwish-KAN, TeLU-KAN, and PWLO-KAN are parameterized as edge-centric linear projection layers within the UKAN architecture.

HardSwish-KAN Linear Layer

HardSwish-KAN integrates the piecewise-linear HardSwish activation into the KAN grid expansion. For an input vector $\mathbf{x} \in \mathbb{R}^{d_{\text{in}}}$, LayerNorm is first applied. A uniform grid of knot points $\{g_0, \dots, g_{G-1}\}$ is constructed across $[g_{\min}, g_{\max}]$. Each feature x_i is evaluated against all grid offsets:

$$\psi_k(x_i) = \text{HardSwish}(x_i - g_k) = (x_i - g_k) \cdot \frac{\text{ReLU6}(x_i - g_k + 3)}{6} \quad (3.7)$$

The evaluations from all G knot points across d_{in} dimensions are concatenated into $\text{vec}(\{\psi_k(x_i)\}) \in \mathbb{R}^{d_{\text{in}} \cdot G}$ and projected via a learnable weight matrix $\mathbf{W} \in \mathbb{R}^{d_{\text{out}} \times (d_{\text{in}} \cdot G)}$.

TeLU-KAN Linear Layer

TeLU-KAN constructs edge-centric transformations by applying the TeLU function over grid-shifted inputs:

$$\psi_k(x_i) = \text{TeLU}(x_i - g_k) = (x_i - g_k) \cdot \tanh(\exp(x_i - g_k)) \quad (3.8)$$

The resulting basis evaluations are concatenated and linearly projected via $\mathbf{W} \in \mathbb{R}^{d_{\text{out}} \times (d_{\text{in}} \cdot G)}$. This leverages TeLU’s identity-like positive region linearity to provide unattenuated gradient propagation across deep KAN blocks.

PWLO-KAN Linear Layer

PWLO-KAN implements an explicit lookup-table (LUT) structure for high-efficiency integer and edge inference. Given domain bounds $[g_{\min}, g_{\max}]$ and segment count G , the interval step size is $\Delta = \frac{g_{\max} - g_{\min}}{G}$. For each input feature x_i , the active segment index k is determined via clamped index quantization:

$$k = \text{clamp} \left(\left\lfloor \frac{x_i - g_{\min}}{\Delta} \right\rfloor, 0, G - 1 \right) \quad (3.9)$$


Rather than evaluating continuous functions, PWLO-KAN retrieves learned affine parameters $a_{i,k}$ and $b_{i,k}$ from embedding tables $\mathbf{A}, \mathbf{B} \in \mathbb{R}^{(d_{\text{in}} \cdot G) \times d_{\text{out}}}$ to evaluate localized affine transformations $\phi_k(x_i) = a_{i,k} \cdot x_i + b_{i,k}$. This enables the entire layer to execute via fast table lookups and linear shift-and-add operations.

3.1.5 Decoder

The decoder combines KAN blocks with standard convolutional upsampling (D_ConvLayer). Each decoder stage: (1) applies a D_ConvLayer convolution to reduce channels, (2) bilinearly upsamples by factor 2, (3) adds the skip connection from the corresponding encoder stage, (4) processes with a decoder KANBlock. The final three stages use only D_ConvLayers to restore the original spatial resolution, followed by a 1×1 convolution head producing per-class logits.

3.2 BiKASegNet: Binarized KAN Segmentation Network

The full BiKA Segmentation Network (we will call it BikaSegNet from now) follows a standard U-Net topology with BiKAConvBlocks replacing standard convolution blocks, as shown in Figure 3.2.



**Note: Red blocks indicate binarized BiKA_Conv2d layers. Blue blocks indicate full-precision FP32 stem and head. Purple block indicates the binary bottleneck. Dashed lines indicate horizontal feature concatenation skip connections across the U-shape.*

Figure 3.2: BiKASegNet Architecture

BiKASegNet implements the KAN philosophy of learnable per-connection activations in a radically different way: instead of smooth spline functions, it uses binarized step functions with learnable per-connection thresholds. This produces a network where both weights and activations are binary ($\{+1, -1\}$), enabling extreme computational efficiency.

3.2.1 Adapting BiKA Layers for Semantic Road Segmentation

While Section 2.3.3 defined the theoretical formulation of the BiKA_Conv2d binarized layer, adapting 1-bit per-connection activation layers for dense semantic road segmentation requires structural and channel-capacity modifications to overcome severe quantization loss and maintain stable gradient propagation across deep networks.

Overall Architectural Adaptation & Channel Width Scaling

To transition BiKA from classification tasks to dense pixel-wise segmentation on high-resolution driving datasets (e.g., BDD100K), an overall architectural adaptation was developed by constructing a U-Net style encoder-decoder network around binarized blocks. In this framework, channel capacity plays a critical role in balancing model expressiveness with hardware efficiency. Two primary channel width scaling configurations were developed and evaluated in the BiKA repository:

- **16-Channel Baseline Configuration** (`base_channels = 16`): In this ultra-lightweight setup, channel dimensions progress across encoder stages as $C \rightarrow 2C \rightarrow 4C \rightarrow 8C$, corresponding to $16 \rightarrow 32 \rightarrow 64 \rightarrow 128$ channels. This model contains only $\sim 0.35\text{M}$ parameters with an extremely small memory footprint ($\sim 1.4\text{ MB}$). It is specifically engineered for sub-watt microcontrollers and embedded FPGAs where memory bandwidth and cache capacities are severely constrained.
- **32-Channel Capacity Update** (`base_channels = 32`): To match the parameter capacity ($\sim 4.1\text{M}$ parameters) and representational bandwidth of full-precision baseline networks (such as UKAN), the network architecture was updated to use a base channel count of $C = 32$. Here, stage channels expand as $32 \rightarrow 64 \rightarrow 128 \rightarrow 256$. Doubling the base channel width quadruples the binary feature interaction capacity

across layers, enabling the network to resolve subtle spatial boundaries, thin lane markings, and distant road obstacles with significantly higher segmentation mIoU, while retaining 100% multiply-free bitwise XNOR and popcount acceleration during inference.

A detailed side-by-side comparison of the 16-channel baseline and the 32-channel capacity update is presented in Table 3.1.

Table 3.1: Side-by-side comparison between 16-channel baseline and 32-channel BiKASegNet configurations.

Architectural Parameter	16-Channel Model ($C = 16$)	32-Channel Model ($C = 32$)
Base Channel Width (C)	16 channels	32 channels
Encoder Feature Maps	$16 \rightarrow 32 \rightarrow 64 \rightarrow 128$	$32 \rightarrow 64 \rightarrow 128 \rightarrow 256$
Decoder Feature Maps	$192 \rightarrow 64 \rightarrow 32 \rightarrow 16$	$384 \rightarrow 128 \rightarrow 64 \rightarrow 32$
Total Model Parameters	$\sim 0.35\text{M}$ parameters	$\sim 4.1\text{M}$ parameters
Memory Footprint	~ 1.4 MB (ultra-lightweight)	~ 4.8 MB ($\sim 32\times$ smaller than 1)
Segmentation Detail	Basic road surface masks	Fine lane boundaries & small ob
Target Deployment	Sub-watt microcontrollers / FPGAs	Edge GPUs (Jetson Nano/Orin)

Full-Precision Input Stem (full_precision_stem)

Across both 16-channel and 32-channel configurations, the initial input stage (`self.enc1`) processes raw RGB images ($3 \times H \times W$) using a full-precision 3×3 convolution ($3 \rightarrow C$), followed by BatchNorm and RReLU, before passing features to binarized BiKA_Conv2d layers. Direct binarization of continuous 8-bit RGB pixel values at the input layer severely degrades spatial details and segmentation accuracy, as raw pixel intensities lack the high-dimensional channel representation required for effective 1-bit thresholding.

ReActNet-Inspired RReLU Activations

Standard ReLU activations set all negative values to zero ($\max(0, x)$), leading to severe information loss and dead neurons in 1-bit binary networks. BiKASegNet replaces standard ReLU with ReActNet-style RReLU (Re-Activation PReLU) across all binary blocks in both 16-channel and 32-

channel variants:

$$\text{RReLU}(x) = \text{ReLU}(x + \gamma) + \beta \quad (3.10)$$

where $\gamma, \beta \in \mathbb{R}^{1 \times C \times 1 \times 1}$ are learnable per-channel pre-shift and post-shift parameters. RReLU dynamically shifts binary activation distributions before and after non-linearities, allowing asymmetric activation ranges necessary for fine-grained pixel-wise road classification.

Full-Precision Classification Head (`full_precision_head`)

For both channel configurations, the final segmentation output layer (`self.final`) uses a 1×1 full-precision convolution projection ($C \rightarrow N_{\text{classes}}$) rather than a binary convolution. Retaining full-precision weights and biases at the final classification head preserves smooth, continuous logit distributions, which are essential for computing precise pixel-wise Cross-Entropy and Dice loss gradients during backpropagation.

Feature Concatenation in U-Net Decoder

In both 16-channel and 32-channel decoders, encoder skip features are concatenated with upsampled decoder features (dec3 takes $8C + 4C = 12C$ channels to $4C$, dec2 takes $4C + 2C = 6C$ channels to $2C$, and dec1 takes $2C + C = 3C$ channels to C). Feature concatenation preserves independent binary feature channels across spatial resolutions, providing richer representation recombination for multi-scale road and lane boundary segmentations.

3.2.2 BiKAConvBlock Design

Each `BiKAConvBlock` integrates two `BiKA_Conv2d` layers with `BatchNorm`, `RReLU` activations, and an FP32 shortcut connection (Bi-Real Net style):

$$\mathbf{s} = \text{Shortcut}(\mathbf{x}) \quad (3.11)$$

$$\mathbf{h} = \text{RReLU}(\text{BN}(\text{BiKA_Conv2d}_1(\mathbf{x}))) \quad (3.12)$$

$$\mathbf{y} = \text{RReLU}(\text{BN}(\text{BiKA_Conv2d}_2(\mathbf{h})) + \mathbf{s}) \quad (3.13)$$

The shortcut `Shortcut( $\mathbf{x}$ )` is a 1×1 FP32 convolution with `BatchNorm` when channel dimensions change, and an identity mapping otherwise. To manage memory consumption during training, activation checkpointing (`torch.utils.checkpoint`) is applied to `BiKAConvBlock`, reducing peak GPU VRAM usage by approximately 60%.

3.3 Optimisation and Training Strategies

While the theoretical formulation of BiKA guarantees multiply-free 1-bit operations, standard PyTorch/ATen operators introduce significant memory transfer bottlenecks when executing 4D per-connection thresholding on GPUs. To maximize both training throughput and inference evaluation efficiency, custom C++/CUDA kernels were engineered in the BiKA repository (`src/bika/csrc/bika_conv2d.cu` and `bika_linear.cu`).

3.3.1 Custom CUDA Kernel Optimisations for BiKA

The custom CUDA kernel pipeline optimizes binary convolution and linear operations through several GPU-level memory and compute strategies:

Bit-Packing & 32-bit Word Quantization

To overcome global memory bandwidth limitations during weight transfers, 32 continuous binary weights along kernel and input channel dimensions are bit-packed into single 32-bit unsigned integers (`unsigned int packed_weight`, i.e., 32-bit computer architecture words). Bit-packing compresses the weight memory footprint by $32\times$, allowing 32 binary synapses to be fetched from GPU VRAM in a single 32-bit memory transaction.

32-Bit Word Alignment and Channel-32 Hardware Optimisation

Because CUDA warp execution processes 32 threads concurrently and hardware popcount (`__popc`) operates on 32-bit unsigned integer words (32-bit `unsigned int` data structures in GPU registers), filter weight dimensions that are exact multiples of 32 eliminate tail remainder masking (`bit < 32`). In `bika_conv2d.cu`, weights are packed into 32-bit words (32-bit integers) along the $C \cdot K_H \cdot K_W$ dimension. When `base_channels = 32`, each 3×3 binary convolution kernel processes $32 \times 3 \times 3 = 288$ elements, which packs into exactly 9 full 32-bit words ($288/32 = 9$ complete 32-bit integers). This eliminates partial bit-masking in thread registers, executing pure full-word bitwise XNOR and `__popc` across every single convolutional stage in the network.

Half-Precision (FP16) Shared Memory Bias Caching

The 4D per-connection bias tensor $\mathbf{b} \in \mathbb{R}^{C_{\text{out}} \times C_{\text{in}} \times K_H \times K_W}$ requires storing an individual threshold per synapse. To prevent VRAM fetch latency, the bias tensor is pre-negated and converted to 16-bit half-precision (`__half neg_bias_half`). At kernel execution, CUDA thread blocks load FP16 bias values directly into fast GPU Shared Memory (`__shared__ __half smem_bias_h`), halving shared memory bandwidth consumption compared to 32-bit float bias loading.

Output-Channel Tiling & Thread-Level Register Reuse

The forward CUDA kernel (`bika_conv2d_forward_kernel`) employs 2D output tiling ($TILE_O = 4$ output channels and $TILE_W = 4$ spatial width pixels per thread). By evaluating 4 output channels concurrently within local thread registers (`acc[TILE_W][TILE_O]`), global memory reads of input feature maps are reduced by $4\times$. Tiling parameters are configured to keep shared memory allocation under 48 KB per block, ensuring high streaming multiprocessor (SM) occupancy.

Bitwise XNOR and Popcount Convolution Execution

During convolution execution, thread registers compare packed 32-bit input activation words (`x_bits`, i.e., 32-bit integer registers) against packed binary weight words (`w_bits`) using bitwise XNOR, followed by the hardware-accelerated population count intrinsic (`__popc`):

$$\text{xnor_val} = \sim (\text{x_bits} \oplus \text{w_bits}) \quad (3.14)$$

$$\text{acc} += 2 \cdot \text{__popc}(\text{xnor_val}) - 32 \quad (3.15)$$

This replaces 32 floating-point multiply-accumulate operations with a single bitwise logic operation and integer popcount, eliminating floating-point ALUs from inner loops.

STE Hard-Tanh Backpropagation Kernels

Custom backward CUDA kernels (`bika_conv2d_backward_wb_kernel` and `bika_linear_backward_kernel`) calculate gradients with respect to inputs ($\frac{\partial \mathcal{L}}{\partial \mathbf{x}}$), weights ($\frac{\partial \mathcal{L}}{\partial \mathbf{W}}$), and per-connection biases ($\frac{\partial \mathcal{L}}{\partial \mathbf{b}}$) using the Straight-Through Estimator (STE). Hard-Tanh gradient clipping ($|z| \leq 1.0$) zeros out gradients for saturated weights, while atomic additions (`atomicAdd`) accumulate weight and bias updates across parallel GPU threads.

Training BiKASegNet requires careful handling of the binary weight dynam-

ics. Specifically:

- **Weight decay exclusion:** Weight decay is explicitly disabled for all BiKA_Conv2d parameters (weights and per-connection biases), BatchNorm affine parameters, and RReLU shift parameters. Applying weight decay to binary thresholds causes them to collapse toward zero, annihilating the STE gradient and killing the backward pass entirely.
- **STE weight clamping:** Binary weights are optionally clamped to $[-1, 1]$ after each optimiser step to ensure they remain within the STE gradient window.
- **Bias initialisation:** Per-connection biases are uniformly initialized in $[-2.0, 0.5]$ to ensure diverse initial activation thresholds across connections.

3.3.2 Data Augmentation

Both UKAN and BiKA training pipelines employ the following augmentations via the Albumentations library:

- Random horizontal flip (probability 0.5).
- Random 90-degree rotation.
- CutMix augmentation for regularisation.
- Resize to the target input resolution (default: 256×256 or 512×512).
- ImageNet-standard normalisation ($\mu = [0.485, 0.456, 0.406]$, $\sigma = [0.229, 0.224, 0.225]$).

Chapter 4

Evaluation and Results

This chapter details the experimental setup, dataset preparation, and performance evaluation of UKAN and BiKASegNet on the BDD100K road segmentation benchmark.

4.1 Experimental Setup

To evaluate full-precision UKAN and binarized BiKASegNet under identical experimental conditions, all models were trained and benchmarked on the BDD100K road segmentation dataset [3]. The experimental pipeline is defined across four structured categories:

- **Hardware Infrastructure Setup:**
 - **Operating System:** Linux
 - **Training GPU:** NVIDIA Quadro RTX 4000 (8 GB VRAM).
 - **Testing / Evaluation GPU:** NVIDIA GeForce RTX 2080 Max-Q (8 GB VRAM).
- **Software Environment Setup:**
 - **Virtual Environment:** Conda environment (`bika-env`).

- **Core Libraries:** Python 3.10+, PyTorch 2.1+ with CUDA 12.1 toolkit, ATen C++ API for custom kernel bindings, Albumentations 1.3+ for spatial/color augmentations, and OpenCV.
- **Training Hyperparameters & Input Dimensions:**
 - **Input shape:** 192x265x3, 360x460x3
 - **Training Epochs:** 50 to 100 epochs trained with AdamW optimiser ($\beta_1 = 0.9, \beta_2 = 0.999$, weight decay = 10^{-4} for FP32 layers and 0.0 for BiKA binary parameters), initial learning rate $\eta_0 = 10^{-3}$, and CosineAnnealingLR scheduler.
 - **Batch Size:** Effective batch size of 16 (batch size 8 per GPU step with gradient accumulation).
- **Output Classes & Segmentation Target Approaches:**
 - **Approach 1: Binary Drivable Road/Lane Segmentation** (`lane_fg / lane2`): $N_{\text{classes}} = 2$ logits ($2 \times H \times W$). Class 0 represents drivable road surface and lane area, while Class 1 represents background. Designed for low-latency autonomous vehicle navigation, collision avoidance, and lane-keeping algorithms.
 - **Approach 2: Multi-Class Full Semantic Segmentation** (`bench19 / drive5`): $N_{\text{classes}} = 19$ logits ($19 \times H \times W$) following the official BDD100K benchmark convention (and $N_{\text{classes}} = 5$ for the 5-class `drive5` collision taxonomy). Evaluates fine-grained multi-class scene understanding across 19 categories (Road, Side-walk, Vehicle, Person, Traffic Light, Sign, Sky, etc.), mapping class index 19 ("unknown") to `ignore_index = 255`.

4.2 Evulation Metrics

The primary evaluation metrics are mean Intersection over Union (mIoU) and mean Dice coefficient, computed via a confusion-matrix-based epoch-level accumulator to avoid batch-averaging inaccuracies.

The IoU for class c is computed as:

$$\text{IoU}_c = \frac{\text{TP}_c}{\text{TP}_c + \text{FP}_c + \text{FN}_c} \quad (4.1)$$

and the Dice coefficient as:

$$\text{Dice}_c = \frac{2 \cdot \text{TP}_c}{2 \cdot \text{TP}_c + \text{FP}_c + \text{FN}_c} \quad (4.2)$$

where TP_c , FP_c , and FN_c are the true positives, false positives, and false negatives for class c accumulated over the entire validation set.

Table 4.1 summarizes the segmentation performance of the evaluated architectures.

Model	Precision	KAN Variant	mIoU (%)	Dice (%)
U-Net (Baseline)	FP32	—	—	—
UKAN (FasterKAN)	FP32	RSWaf	—	—
UKAN (ReLU-KAN)	FP32	ReLU	—	—
UKAN (PWLO-KAN)	FP32	PWLO	—	—
BiKASegNet	1-bit	Binary Step	—	—
BiKASegNet + KD	1-bit	Binary Step	—	—

**Note: Results to be filled from experimental runs. KD = Knowledge Distillation from UKAN teacher.*

Table 4.1: Segmentation accuracy comparison across architectures on BDD100K road segmentation.

4.3 Results and Deiscussion

Table 4.2 compares the model size and computational efficiency of the evaluated architectures.

The theoretical compression ratio of BiKASegNet relative to a full-precision network of identical topology is approximately $32\times$ for the binarized layers, since each 32-bit float weight is replaced by a single bit. In practice, the full-precision stem and head slightly reduce the overall compression ratio.

Model	Params (M)	Size (MB)	FPS	Compression
U-Net (FP32)	—	—	—	$1\times$
UKAN (FP32)	—	—	—	$1\times$
BiKASegNet (1-bit)	—	—	—	$\sim 32\times$

**Note: Compression ratio is relative to the FP32 model of equivalent architecture. FPS measured on the evaluation GPU.*

Table 4.2: Model efficiency comparison: parameter count, storage size, inference throughput, and compression ratio.

4.3.1 Inference Throughput

The binary convolution operations in BiKA_Conv2d leverage XNOR and popcount operations through the custom CUDA kernel. Table 4.3 profiles the per-layer latency contributions.

Component	UKAN Latency (ms)	BiKA Latency (ms)
Encoder Stages	—	—
KAN/Binary Blocks	—	—
Decoder Stages	—	—
Classification Head	—	—
Total	—	—

**Note: Latency measured as mean over 1000 forward passes at 512×512 input resolution.*

Table 4.3: Per-component inference latency comparison between UKAN and BiKASegNet.

To isolate the impact of the KAN activation function choice, we trained UKAN with each variant using identical hyperparameters. The KAN variant selection affects both accuracy and inference speed, as each variant trades off the smoothness of the learned activation against computational cost.

Table 4.4 examines the impact of knowledge distillation on BiKASegNet accuracy.

We evaluate the individual contribution of each BNN accuracy-improvement technique integrated into BiKAConvBlock:

Configuration	Teacher	mIoU (%)	Δ mIoU
BiKA (no KD)	—	—	—
BiKA + KD ($\tau = 4$)	UKAN-FasterKAN	—	—
BiKA + KD + Boundary	UKAN-FasterKAN	—	—

**Note: Δ mIoU is relative to the no-KD baseline.*

Table 4.4: Ablation study on knowledge distillation and boundary loss for BiKASegNet.

- **Libra-PB**: Weight centering for balanced binarization.
- **RPreLU**: Learnable activation shifts vs. standard ReLU.
- **FP32 Skip**: Bi-Real Net gradient highway connections.
- **AdaBin**: Learnable output scaling.

Disabling the legacy mode activates all four techniques simultaneously. In legacy mode, the network reverts to plain ReLU activations and no skip connections, establishing the minimal BNN baseline.

Chapter 5

Conclusions and Future Works

5.1 Conclusions

In this thesis, we implemented, evaluated, and compared two complementary approaches for integrating Kolmogorov-Arnold Networks into segmentation architectures for road scene understanding.

The UKAN architecture demonstrates that replacing standard MLP projections with KAN layers in the U-Net bottleneck produces competitive segmentation quality while offering greater representational flexibility through learnable per-edge activation functions. The availability of multiple KAN variants (FasterKAN, ReLU-KAN, HardSwish-KAN, PWLO-KAN, TeLU-KAN) enables architecture-level tuning of the accuracy-latency trade-off depending on the deployment target.

The BiKASegNet architecture demonstrates that the KAN philosophy of per-connection learnable activations can be realized through binarized step functions with learnable thresholds, achieving extreme compression ($\sim 32\times$) and computational acceleration through XNOR/popcount operations. The integration of state-of-the-art BNN techniques (Libra-PB, AdaBin, RPreLU, FP32 skip connections) is essential for maintaining segmentation quality un-

der binary constraints. Knowledge distillation from a full-precision UKAN teacher further bridges the binarization gap.

5.2 Future Works

To extend this research, several areas should be investigated:

- **Edge Deployment:** Deploy BiKASegNet on actual embedded platforms (Raspberry Pi, Jetson Nano) and benchmark real-world inference latency and power consumption.
- **Multi-class Segmentation:** Extend the evaluation from 2-class road segmentation to the full BDD100K label set (20+ classes) and evaluate how KAN integration affects fine-grained category discrimination.
- **Temporal Consistency:** Integrate the segmentation networks into a video pipeline with temporal smoothing to evaluate frame-to-frame consistency of road boundary predictions.
- **Hybrid KAN-BNN Architectures:** Explore architectures that selectively apply binarization only to specific encoder/decoder stages while keeping KAN-critical bottleneck layers in higher precision.

Bibliography

- [1] Z. Liu, Y. Wang, S. Vaidya, F. Ruehle, J. Halverson, M. Soljačić, T. Y. Hou, and M. Tegmark, “KAN: Kolmogorov-arnold networks,” *arXiv preprint arXiv:2404.19756*, 2024.
- [2] O. Ronneberger, P. Fischer, and T. Brox, “U-Net: Convolutional networks for biomedical image segmentation,” in *International Conference on Medical Image Computing and Computer-Assisted Intervention (MICCAI)*, pp. 234–241, 2015.
- [3] F. Yu, H. Chen, X. Wang, W. Xian, Y. Chen, F. Liu, V. Madhavan, and T. Darrell, “BDD100K: A diverse driving dataset for heterogeneous multitask learning,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 2636–2645, 2020.
- [4] Y. Liu, S. Ullah, and A. Kumar, “BiKA: Kolmogorov-Arnold-Network-inspired ultra lightweight neural network hardware accelerator,” *arXiv preprint arXiv:2602.23455*, 2026.
- [5] J. Long, E. Shelhamer, and T. Darrell, “Fully convolutional networks for semantic segmentation,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 3431–3440, 2015.
- [6] E. Xie, W. Wang, Z. Yu, A. Anandkumar, J. M. Alvarez, and P. Luo, “SegFormer: Simple and efficient design for semantic segmentation with transformers,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 34, pp. 12077–12090, 2021.

- [7] H. Cao, Y. Wang, J. Chen, D. Jiang, X. Zhang, Q. Tian, and M. Wang, “Swin-Unet: Unet-like pure transformer for medical image segmentation,” *arXiv preprint arXiv:2105.05537*, 2022.
- [8] A. Kirillov, E. Mintun, N. Ravi, H. Mao, C. Rolland, L. Gustafson, T. Xiao, S. Whitehead, A. C. Berg, W.-Y. Lo, P. Dollár, and R. Girshick, “Segment anything,” *arXiv preprint arXiv:2304.02643*, 2023.
- [9] N. Silberman, D. Hoiem, P. Kohli, and R. Fergus, “Indoor segmentation and support inference from RGBD images,” in *Proceedings of the European Conference on Computer Vision (ECCV)*, pp. 746–760, 2012.
- [10] V. Badrinarayanan, A. Kendall, and R. Cipolla, “SegNet: A deep convolutional encoder-decoder architecture for image segmentation,” *IEEE Transactions on Pattern Analysis and Machine Intelligence*, vol. 39, no. 12, pp. 2481–2495, 2017.
- [11] L.-C. Chen, Y. Zhu, G. Papandreou, F. Schroff, and H. Adam, “Encoder-decoder with atrous separable convolution for semantic image segmentation,” in *Proceedings of the European Conference on Computer Vision (ECCV)*, pp. 801–818, 2018.
- [12] H. Zhao, J. Shi, X. Qi, X. Wang, and J. Jia, “Pyramid scene parsing network,” in *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 2881–2890, 2017.
- [13] A. Paszke, A. Chaurasia, S. Kim, and E. Culurciello, “ENet: A deep neural network architecture for real-time semantic segmentation,” *arXiv preprint arXiv:1606.02147*, 2016.
- [14] C. Yu, J. Wang, P. C. Chao, C. Gao, N. Sang, and J. Liu, “BiSeNet: Bilateral segmentation network for real-time semantic segmentation,” in *Proceedings of the European Conference on Computer Vision (ECCV)*, pp. 325–341, 2018.
- [15] A. N. Kolmogorov, “On the representation of continuous functions of several variables by superposition of continuous functions of one variable

- and addition,” *Doklady Akademii Nauk SSSR*, vol. 114, pp. 953–956, 1957.
- [16] C. Li, X. Liu, W. Li, C. Wang, H. Liu, and Y. Yuan, “U-KAN makes strong backbone for medical image segmentation and generation,” *arXiv preprint arXiv:2406.02918*, 2024.
- [17] A. Tsiamis, “FasterKAN: Accelerating kolmogorov-arnold networks with RSWAf basis functions.” <https://github.com/AthanasiosDelis/faster-kan>, 2024.
- [18] I. Hubara, M. Courbariaux, D. Soudry, R. El-Yaniv, and Y. Bengio, “Binarized neural networks: Training deep neural networks with weights and activations constrained to +1 or -1,” in *Advances in Neural Information Processing Systems (NeurIPS)*, vol. 29, pp. 4107–4115, 2016.
- [19] M. Rastegari, V. Ordonez, J. Redmon, and A. Farhadi, “XNOR-Net: Imagenet classification using binary convolutional neural networks,” in *Proceedings of the European Conference on Computer Vision (ECCV)*, pp. 525–542, 2016.
- [20] Z. Liu, B. Wu, W. Luo, X. Yang, W. Liu, and K.-T. Cheng, “Bi-Real Net: Enhancing the performance of 1-bit CNNs with improved representational capability and advanced training algorithm,” in *Proceedings of the European Conference on Computer Vision (ECCV)*, pp. 722–737, 2018.
- [21] Z. Liu, Z. Shen, M. Savvides, and K.-T. Cheng, “ReActNet: Towards precise binary neural network with generalized activation functions,” in *Proceedings of the European Conference on Computer Vision (ECCV)*, pp. 143–159, 2020.
- [22] H. Qin, R. Gong, X. Liu, M. Shen, Z. Wei, F. Yu, and J. Song, “Forward and backward information retention for accurate binary neural networks,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 2250–2259, 2020.

- [23] Z. Tu, X. Chen, P. Ren, and Y. Wang, “AdaBin: Improving binary neural networks with adaptive binary sets,” in *Proceedings of the European Conference on Computer Vision (ECCV)*, pp. 262–278, 2022.
- [24] B. Zhuang, C. Shen, M. Tan, L. Liu, and I. Reid, “Structured binary neural networks for accurate image classification and semantic segmentation,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 413–422, 2019.
- [25] V. I. Arnold, “On the representation of functions of several variables as a superposition of functions of a smaller number of variables,” *Selected Works*, 2009.
- [26] Z. Li, “Kolmogorov-arnold networks are radial basis function networks,” *arXiv preprint arXiv:2405.06721*, 2024.
- [27] R. Qiu, Y. Miao, S. Wang, L. Yu, Y. Zhu, and X.-S. Gao, “PowerMLP: An efficient version of KAN,” *arXiv preprint arXiv:2412.13571*, 2024.
- [28] Q. Qiu, G. Xu, *et al.*, “ReLU-KAN: New kolmogorov-arnold networks that only need matrix addition, dot multiplication, and ReLU,” *arXiv preprint arXiv:2406.02075*, 2024.
- [29] A. Howard, M. Sandler, G. Chu, L.-C. Chen, B. Chen, M. Tan, W. Wang, Y. Zhu, R. Pang, V. Vasudevan, Q. V. Le, and H. Adam, “Searching for MobileNetV3,” in *Proceedings of the IEEE/CVF International Conference on Computer Vision (ICCV)*, pp. 1314–1324, 2019.
- [30] A. Fernandez and A. Mali, “TeLU activation function for fast and stable deep learning,” *arXiv preprint arXiv:2412.20269*, 2024.
- [31] N. Schoots, M. J. Villani, and N. uit de Bos, “Relating piecewise linear kolmogorov arnold networks to ReLU networks,” in *Proceedings of the 28th International Conference on Artificial Intelligence and Statistics (AISTATS)*, vol. 258, pp. 199–207, 2025.

-
- [32] Y. Bengio, N. Léonard, and A. Courville, “Estimating or propagating gradients through stochastic neurons for conditional computation,” in *arXiv preprint arXiv:1308.3432*, 2013.